

**SUBMITTED BY: MONIKA SINGH**

## **Blogs /Articles Details**

### **Problem Statement:**

Every year a lot of companies hire a number of employees. The companies invest time and money in training those employees, not just this but there are training programs within the companies for their existing employees as well. The aim of these programs is to increase the effectiveness of their employees. But where HR Analytics fit in this? and is it just about improving the performance of employees?

### **HR Analytics:**

Human resource analytics (HR analytics) is an area in the field of analytics that refers to applying analytic processes to the human resource department of an organization in the hope of improving employee performance and therefore getting a better return on investment. HR analytics does not just deal with gathering data on employee efficiency. Instead, **it aims to provide insight into each process by gathering data and then using it to make relevant decisions about how to improve these processes.**

### **Attrition in HR:**

Attrition in human resources refers to the gradual loss of employees overtime. In general, relatively high attrition is problematic for companies. HR professionals often assume a leadership role in designing company compensation programs, work culture, and motivation systems that help the organization retain top employees.

How does Attrition affect companies? and how does HR Analytics help in analyzing attrition? We will discuss the first question here and for the second question, we will write the code and try to understand the process step by step.

### **Attrition affecting Companies:**

A major problem in high employee attrition is its cost to an organization. Job postings, hiring processes, paperwork, and new hire training are some of the common expenses of losing employees and replacing them. Additionally, regular employee turnover prohibits your organization from increasing its collective knowledge base and experience over time. This is especially concerning if your business is customer-facing, as customers often prefer to interact with familiar people. Errors and issues are more likely if you constantly have new workers.

# Importing Libraries

```
In [10]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [11]: df=pd.read_csv(r'HR-Employee-Attrition.csv')
df
```

Out[11]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	...	RelationshipSatisfacti
--	-----	-----------	----------------	-----------	------------	------------------	-----------	----------------	---------------	----------------	-----	------------------------

0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	...	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	...	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	...	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	...	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	...	
...	...	...	...	...	...	...	...	...	...	...	...	
1465	36	No	Travel_Frequently	884	Research & Development	23	2	Medical	1	2061	...	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	Medical	1	2062	...	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	Life Sciences	1	2064	...	
1468	49	No	Travel_Frequently	1023	Sales	2	3	Medical	1	2065	...	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	Medical	1	2068	...	

1470 rows × 35 columns



```
In [12]: #checking the shape of the data  
df.shape
```

```
Out[12]: (1470, 35)
```

```
In [13]: #checking the info of the data set  
df.info()
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EmployeeCount	1470 non-null	int64
9	EmployeeNumber	1470 non-null	int64
10	EnvironmentSatisfaction	1470 non-null	int64
11	Gender	1470 non-null	object
12	HourlyRate	1470 non-null	int64
13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object

```

RelationshipSatisfaction    int64
StandardHours              int64
StockOptionLevel           int64
TotalWorkingYears          int64
TrainingTimesLastYear      int64
WorkLifeBalance            int64
YearsAtCompany             int64
YearsInCurrentRole         int64
YearsSinceLastPromotion    int64
YearsWithCurrManager       int64
dtype: object

```

In this dataset we have 9 attributes of object data types and 27 attributes of integer.

```

In [15]: #checking the null values
df.isnull().sum()

```

```

Out[15]: Age                0
Attrition                  0
BusinessTravel            0
DailyRate                 0
Department                0
DistanceFromHome          0
Education                 0
EducationField            0
EmployeeCount             0
EmployeeNumber            0
EnvironmentSatisfaction   0
Gender                    0
HourlyRate                0
JobInvolvement            0
JobLevel                  0
JobRole                   0
JobSatisfaction           0

```

Activate Windows  
Go to Settings to activate Windows

```
In [14]: #checking the data types
df.dtypes
```

```
Out[14]: Age                int64
Attrition                  object
BusinessTravel             object
DailyRate                 int64
Department                 object
DistanceFromHome          int64
Education                 int64
EducationField             object
EmployeeCount             int64
EmployeeNumber            int64
EnvironmentSatisfaction   int64
Gender                    object
HourlyRate                int64
JobInvolvement            int64
JobLevel                  int64
JobRole                   object
JobSatisfaction           int64
MaritalStatus             object
MonthlyIncome             int64
MonthlyRate               int64
NumCompaniesWorked        int64
Over18                    object
OverTime                  object
PercentSalaryHike         int64
PerformanceRating         int64
RelationshipSatisfaction   int64
StandardHours             int64
```

```
In [15]: #checking the null values
df.isnull().sum()
```

```
Out[15]: Age                                0
Attrition                                  0
BusinessTravel                             0
DailyRate                                  0
Department                                 0
DistanceFromHome                           0
Education                                  0
EducationField                              0
EmployeeCount                              0
EmployeeNumber                             0
EnvironmentSatisfaction                    0
Gender                                      0
HourlyRate                                 0
JobInvolvement                             0
JobLevel                                   0
JobRole                                    0
JobSatisfaction                            0
MaritalStatus                             0
MonthlyIncome                             0
MonthlyRate                                0
NumCompaniesWorked                         0
Over18                                     0
OverTime                                    0
PercentSalaryHike                          0
PerformanceRating                         0
RelationshipSatisfaction                   0
StandardHours                             0
StockOptionLevel                          0
```

# Heatmap

```
In [16]: #plotting the heatmap for checking the null value
plt.figure(figsize=(15,10))
sns.heatmap(df.isnull());
```




```
In [17]: #stats
df.describe()
```

```
Out[17]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	...	Rela
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	...	
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.891156	2.729932	2.063946	...	
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.329428	0.711561	1.106940	...	
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.000000	1.000000	1.000000	...	
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.000000	2.000000	1.000000	...	
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.000000	3.000000	2.000000	...	
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.750000	3.000000	3.000000	...	
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.000000	4.000000	5.000000	...	

8 rows × 26 columns



```
In [18]: #checking the unique values in column
print(df['EmployeeCount'].unique())
print(df['StandardHours'].unique())
```

```
[1]
[80]
```

```
In [20]: # dropping duplicate cols
df_1 = df.drop(['EmployeeCount', 'StandardHours', 'EmployeeNumber'],axis=1)
```

## Categorical Attributes

```
In [21]: # segregating the object datatype.
ob=df_1.select_dtypes(include='object')
```

```
In [22]: ob
```

```
Out[22]:
```

	Attrition	BusinessTravel	Department	EducationField	Gender	JobRole	MaritalStatus	Over18	OverTime
0	Yes	Travel_Rarely	Sales	Life Sciences	Female	Sales Executive	Single	Y	Yes
1	No	Travel_Frequently	Research & Development	Life Sciences	Male	Research Scientist	Married	Y	No
2	Yes	Travel_Rarely	Research & Development	Other	Male	Laboratory Technician	Single	Y	Yes
3	No	Travel_Frequently	Research & Development	Life Sciences	Female	Research Scientist	Married	Y	Yes
4	No	Travel_Rarely	Research & Development	Medical	Male	Laboratory Technician	Married	Y	No
...	...	...	...	...	...	...	...	...	...
1465	No	Travel_Frequently	Research & Development	Medical	Male	Laboratory Technician	Married	Y	No
1466	No	Travel_Rarely	Research & Development	Medical	Male	Healthcare Representative	Married	Y	No
1467	No	Travel_Rarely	Research & Development	Life Sciences	Male	Manufacturing Director	Married	Y	Yes
1468	No	Travel_Frequently	Sales	Medical	Male	Sales Executive	Married	Y	No
1469	No	Travel_Rarely	Research & Development	Medical	Male	Laboratory Technician	Married	Y	No


1470 rows × 9 columns

Activate W  
Go to Setting:



## Data Visualization

```
In [24]: plt.figure(figsize=(10,8))
plt.title('Target Data Distribution',fontsize=16)
sns.countplot(df_1['Attrition'],data=df_1);
```



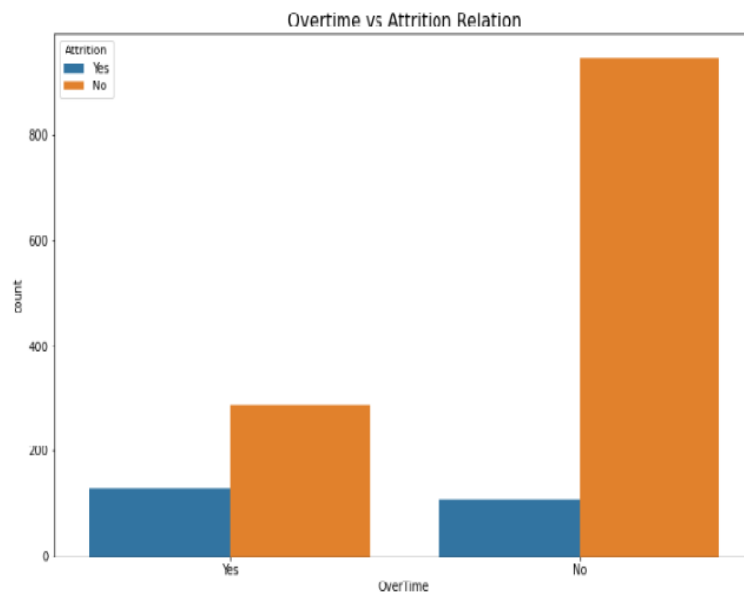
Data is imbalanced in nature.

Activate Windows  
Go to Settings to activate

```
In [25]:
```

Data is imbalanced in nature.

```
In [25]: plt.figure(figsize=(12,8))
plt.title('Overtime vs Attrition Relation',fontsize=15)
sns.countplot(df_1['OverTime'],hue='Attrition',data=df_1);
```




People who do over time have higher chances to left the company as compair to person who don't do overtime.

Activate Window  
Go to Settings to activate

```
In [26]: plt.figure(figsize=(12,8))
```

People who do overtime have higher chances to leave the company as compared to people who don't do overtime.

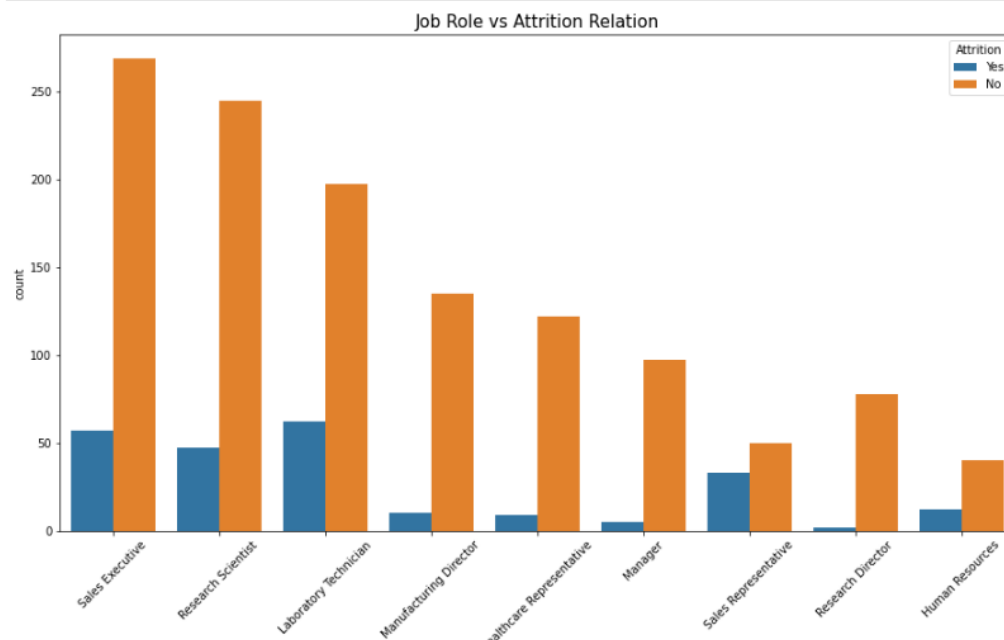
```
In [26]: plt.figure(figsize=(12,8))
plt.title('MaritalStatus vs Attrition Relation',fontsize=15)
sns.countplot(df_1['MaritalStatus'],hue='Attrition',data=df_1);
```



unmarried people have higher tendency to leave the company as compared to Married

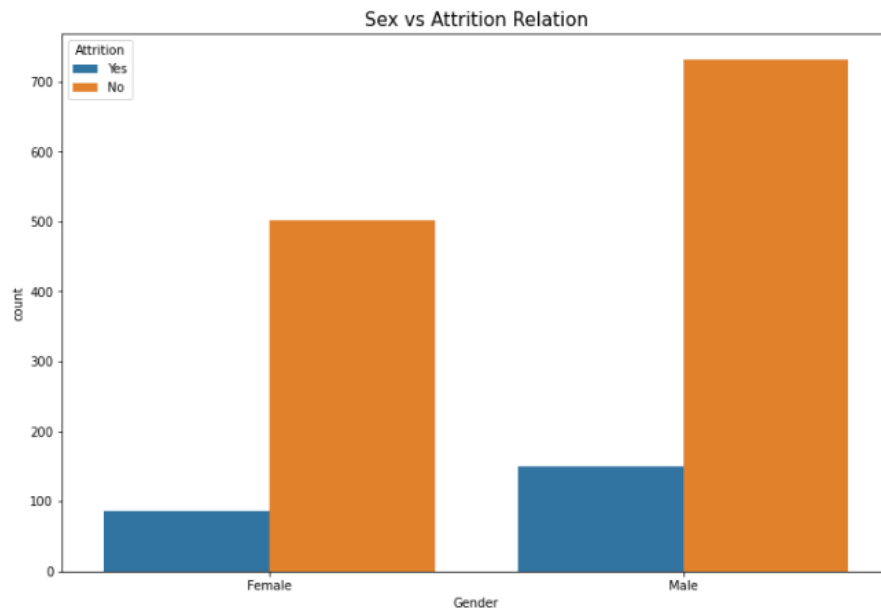
Divorced/separated has lowest chances to leave the company.

```
In [27]: plt.figure(figsize=(15,8))
plt.xticks(rotation=45)
plt.title('Job Role vs Attrition Relation',fontsize=15)
sns.countplot(df_1['JobRole'],hue='Attrition',data=df_1);
```



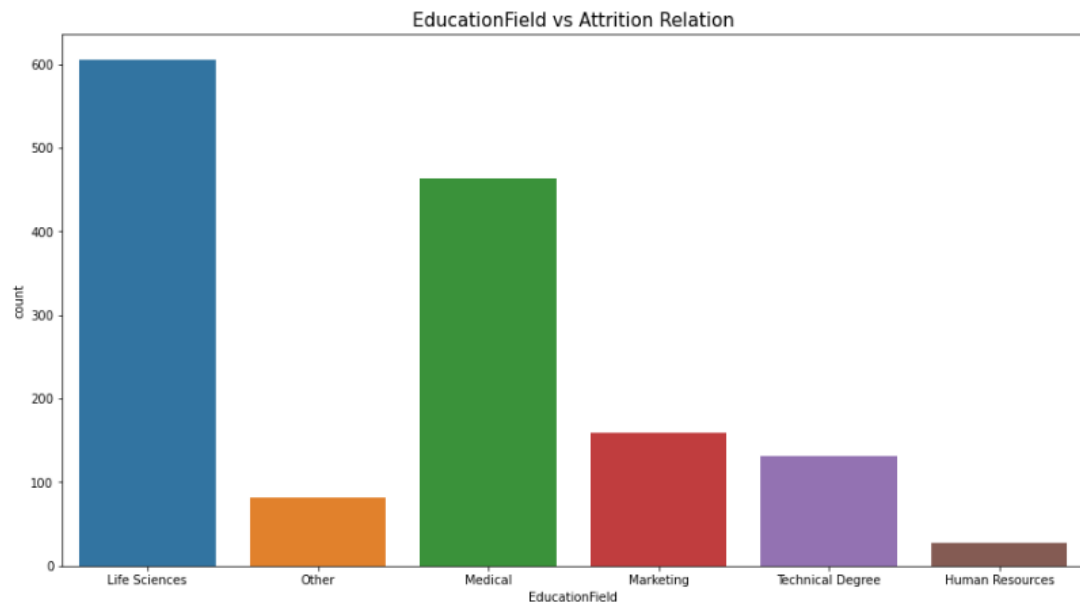
Laboratory Technician ,Sales Executive, Sales representatives research scientis have higher tendency to leave job

```
In [28]: plt.figure(figsize=(12,8))
plt.title('Sex vs Attrition Relation',fontsize=15)
sns.countplot(df_1['Gender'],hue='Attrition',data=df_1);
```



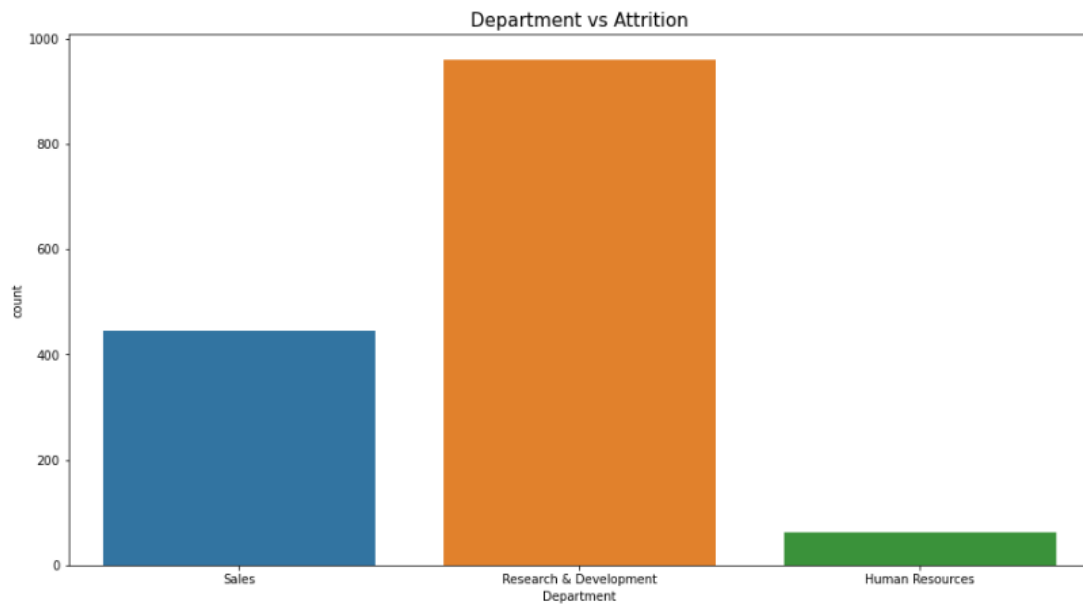
```
In [29]: plt.figure(figsize=(10,8))
plt.title('Sex vs Attrition',fontsize=15)
sns.countplot(df_1['Gender'],data=df_1[df_1['Attrition']=='Yes']);
```

```
In [30]: plt.figure(figsize=(15,8))
plt.title('EducationField vs Attrition Relation',fontsize=15)
sns.countplot(df_1['EducationField'],data=df_1[df_1['Attrition']=='Yes']);
```



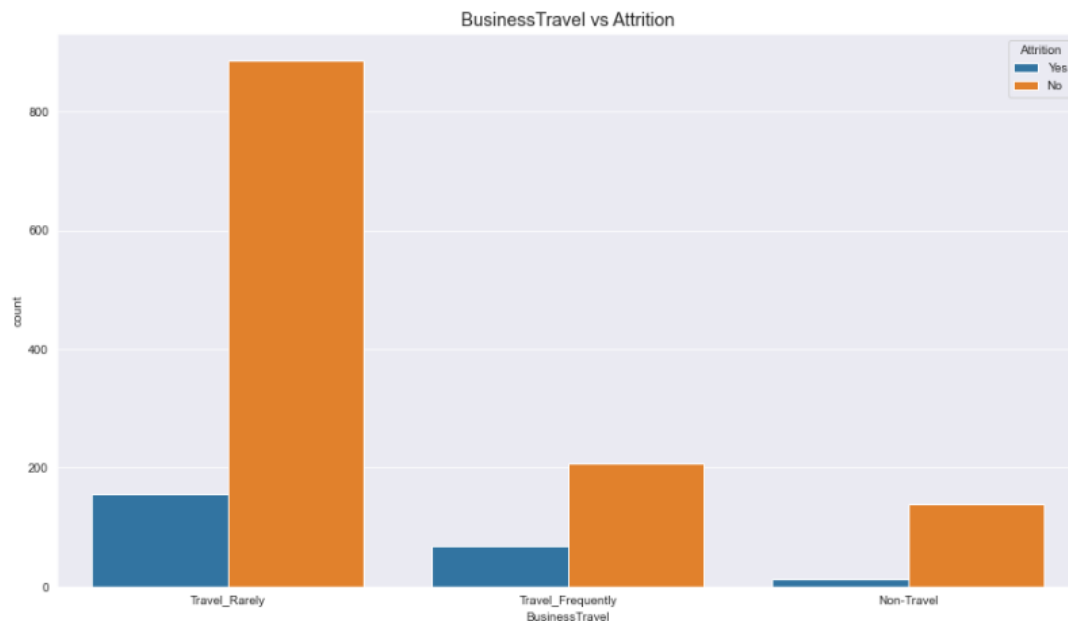
life science and Medical two major Education fields which has higher demand in market and people are switch their job more frequent.

```
In [31]: plt.figure(figsize=(15,8))
plt.title('Department vs Attrition',fontsize=15)
sns.countplot(df_1['Department'],data=df_1[df_1['Attrition']=='Yes']);
```




people who work in Research & development have higher chances to leave the company.

```
In [32]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('BusinessTravel vs Attrition',fontsize=15)
sns.countplot(df_1['BusinessTravel'],hue='Attrition',data=df_1);
```



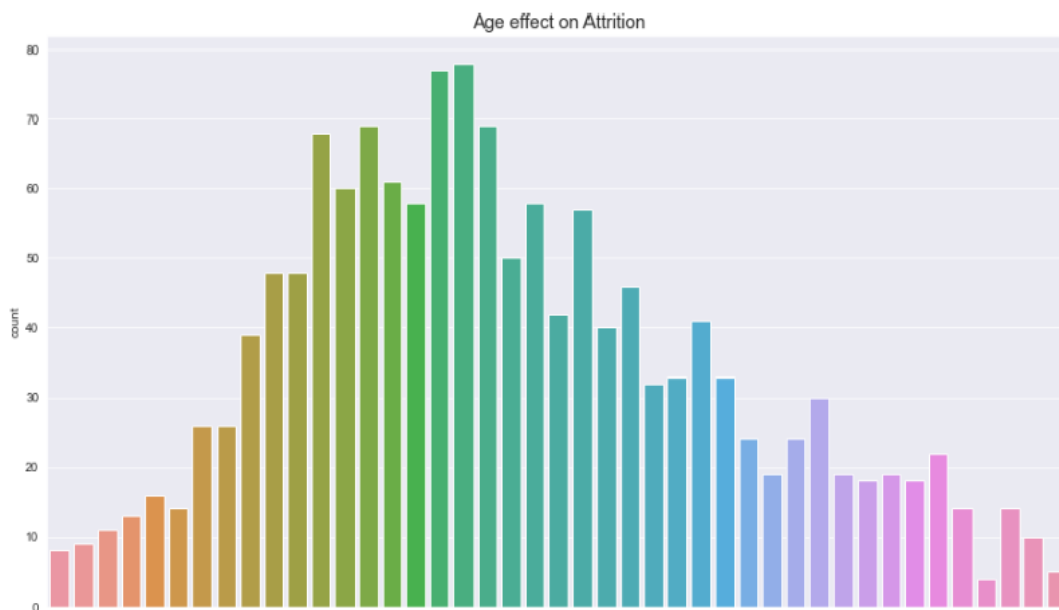
```
In [33]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('BusinessTravel vs Attrition',fontsize=15)
sns.countplot(df_1['BusinessTravel'],data=df_1[df_1['Attrition']=='Yes']);
```



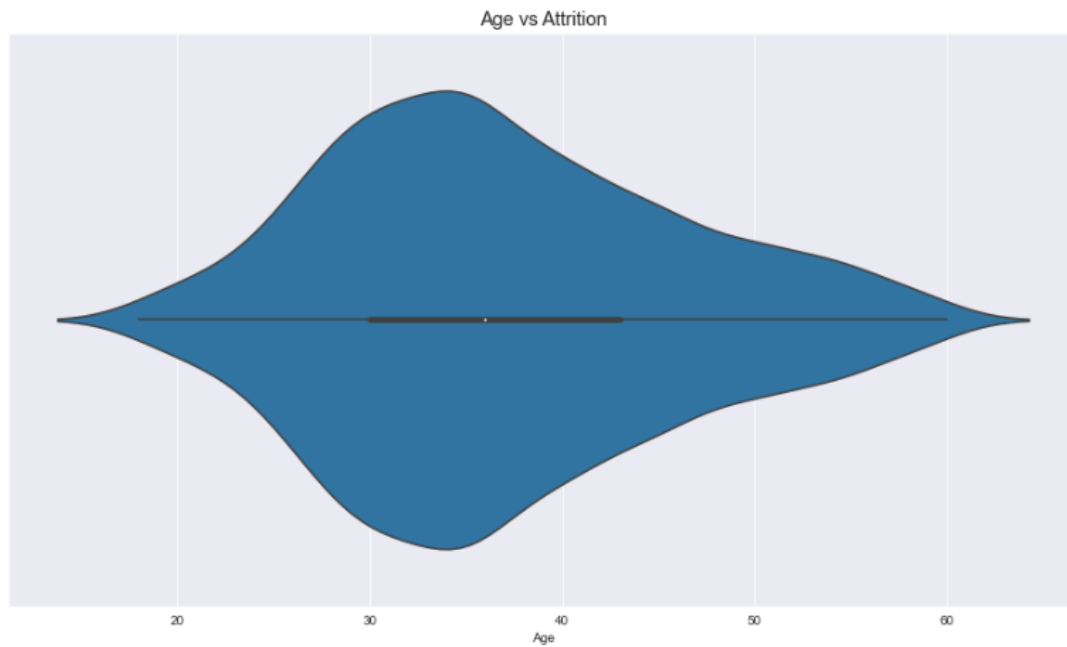
Both countplot it is clear that an employee who travels rarely to other places have higher chances to left job.

Both countplot it is clear that an employee who travels rarely to other places have higher chances to left job.

```
In [34]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Age effect on Attrition',fontsize=15)
sns.countplot(df_1['Age'],data=df_1[df_1['Attrition']=='Yes']);
```

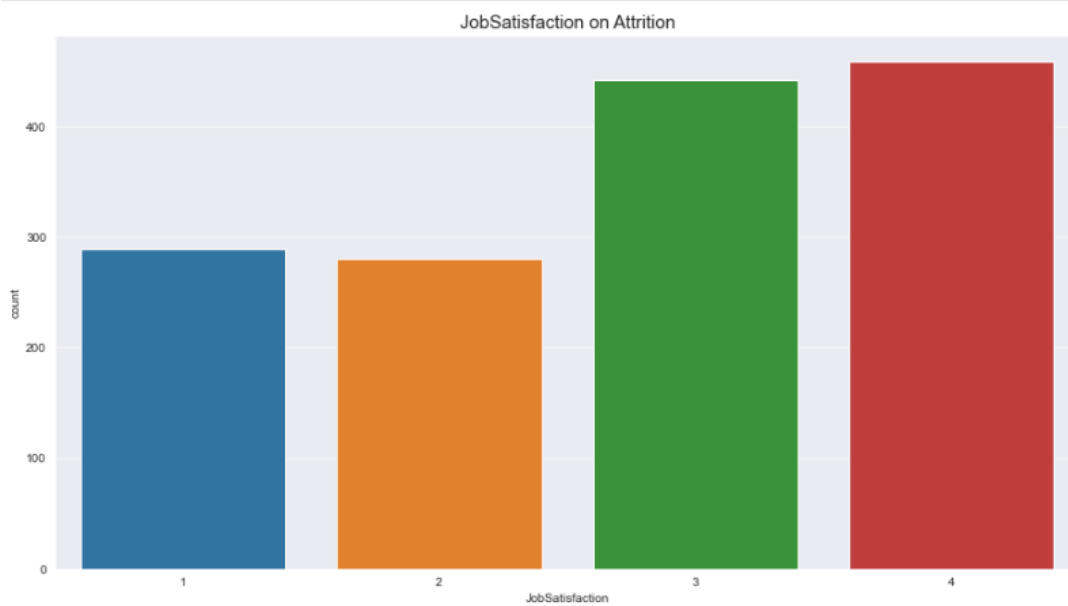


```
In [35]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Age vs Attrition',fontsize=15)
sns.violinplot(df_1['Age'],data=df_1[df_1['Attrition']=='Yes']);
```

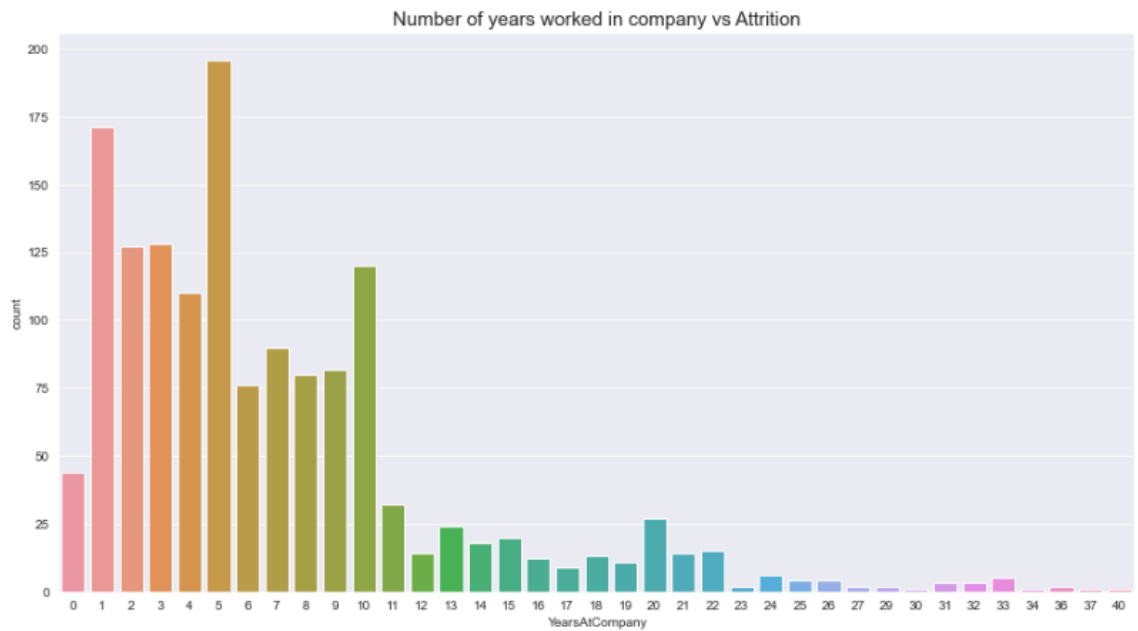


Two plots we can visualize that a employees whose age group is between 28-39 has highest tendency to job change


```
In [36]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('JobSatisfaction on Attrition',fontsize=15)
sns.countplot(df_1['JobSatisfaction'],data=df_1[df_1['Attrition']=='Yes']);
```



```
In [37]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Number of years worked in company vs Attrition',fontsize=15)
sns.countplot(df_1['YearsAtCompany'],data=df_1[df_1['Attrition']=='Yes']);
```



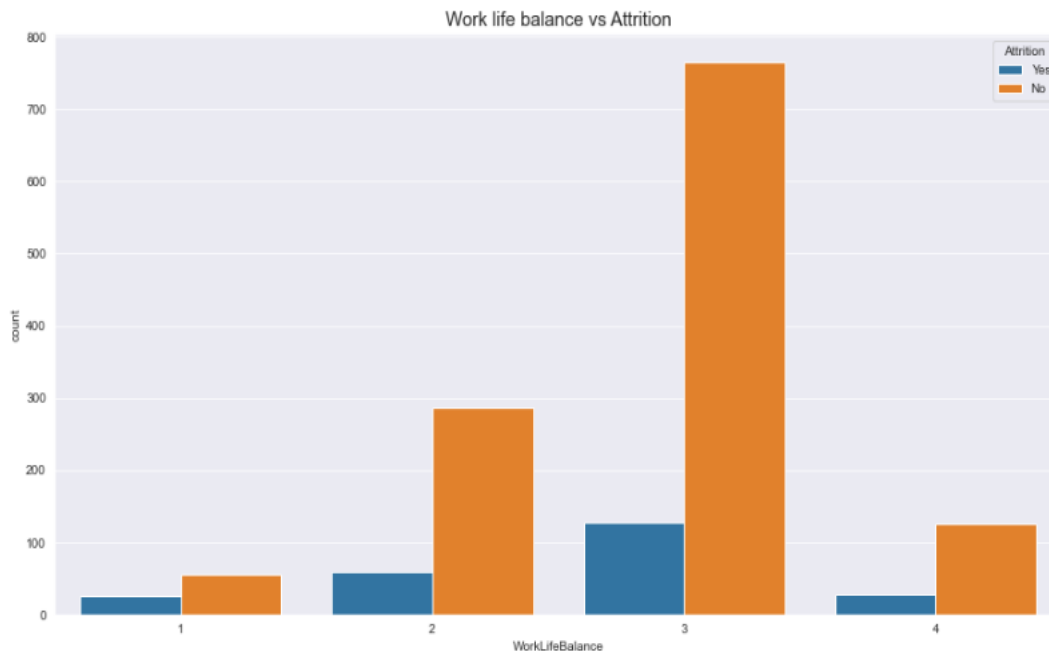
```
In [38]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Last Promotion vs Attrition',fontsize=15)
sns.countplot(df_1['YearsSinceLastPromotion'],hue='Attrition',data=df_1);
```



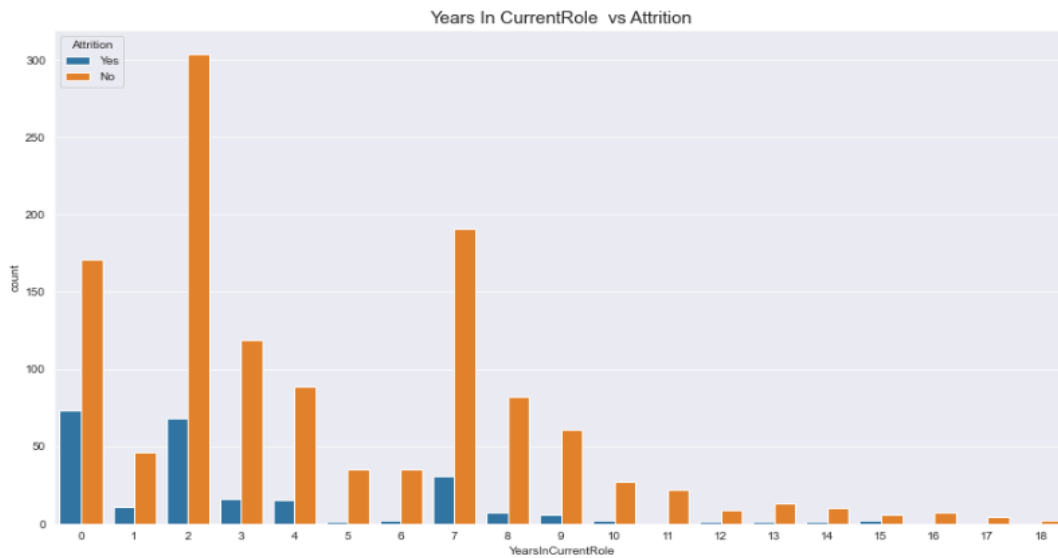
who either got promoted recently or Joined recently has the highest chances of job change. Employee who did not get promotion for 1-2 years also hav for job change.



```
In [40]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Work life balance vs Attrition',fontsize=15)
sns.countplot(df_1['WorkLifeBalance'],hue='Attrition',data=df_1);
```

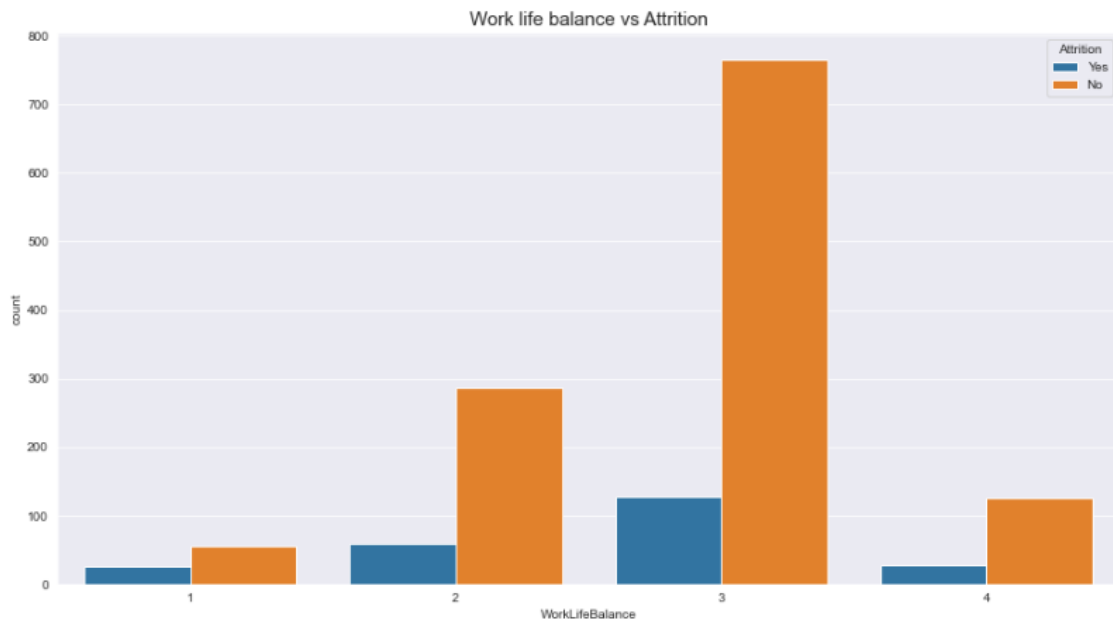


```
In [39]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Years In CurrentRole vs Attrition',fontsize=15)
sns.countplot(df_1['YearsInCurrentRole'],hue='Attrition',data=df_1);
```

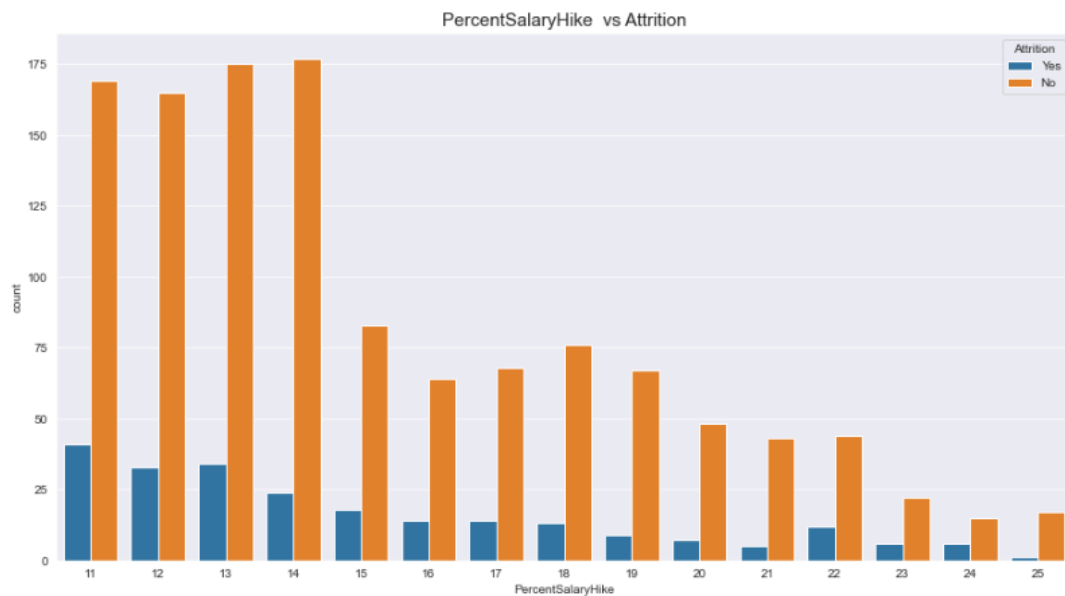


person who has experience of 2 years or who joined recently and working on same profile has fair chances to left job

```
In [40]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('Work life balance vs Attrition',fontsize=15)
sns.countplot(df_1['WorkLifeBalance'],hue='Attrition',data=df_1);
```



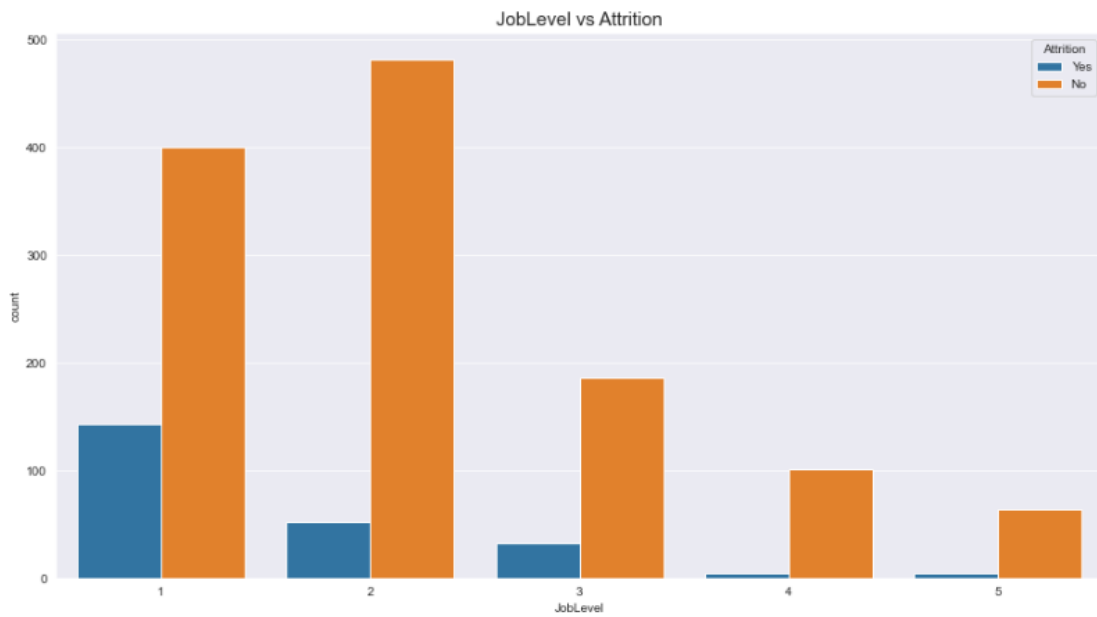
```
In [41]: plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('PercentSalaryHike vs Attrition',fontsize=15)
sns.countplot(df_1['PercentSalaryHike'],hue='Attrition',data=df_1);
```



It is clearly visible that an employee who got least salary hike in year have higher tendency for job change.

In [42]:


```
plt.figure(figsize=(15,8))
sns.set_style('darkgrid')
plt.title('JobLevel vs Attrition',fontsize=15)
sns.countplot(df_1['JobLevel'],hue='Attrition',data=df_1);
```



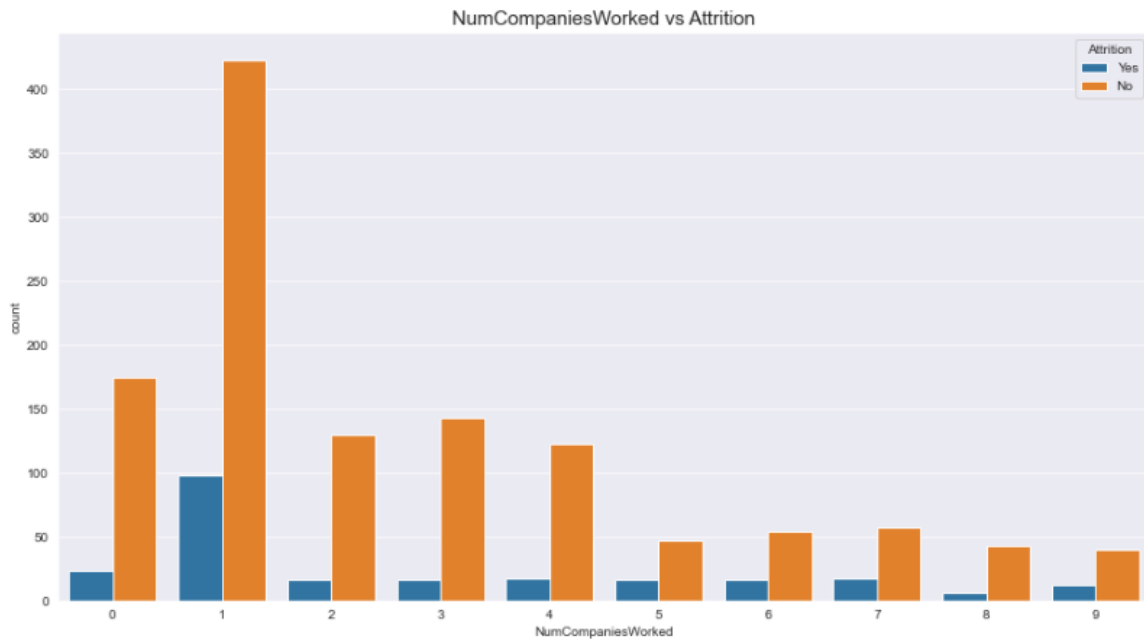
Employee who has job band/level 1 has tendency to change job more frequent then job level 4 or 5

In [43]:

```
plt.figure(figsize=(15,8))
plt.title('Monthly Income vs Attrition',fontsize=15)
sns.stripplot(df_1['Attrition'],df_1['MonthlyIncome'],data=df_1);
```



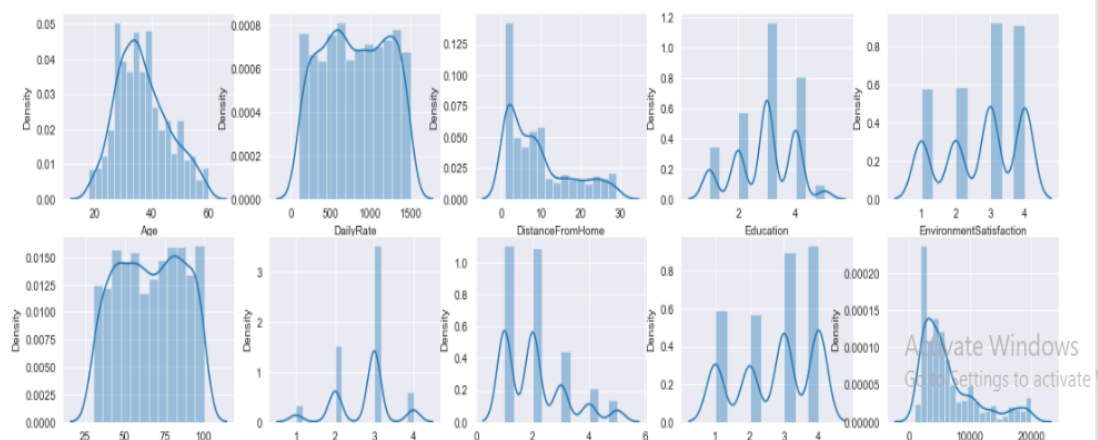
```
[n [44]: plt.figure(figsize=(15,8))
plt.title('NumCompaniesWorked vs Attrition',fontsize=15)
sns.countplot(df_1['NumCompaniesWorked'],hue='Attrition',data=df_1);
```

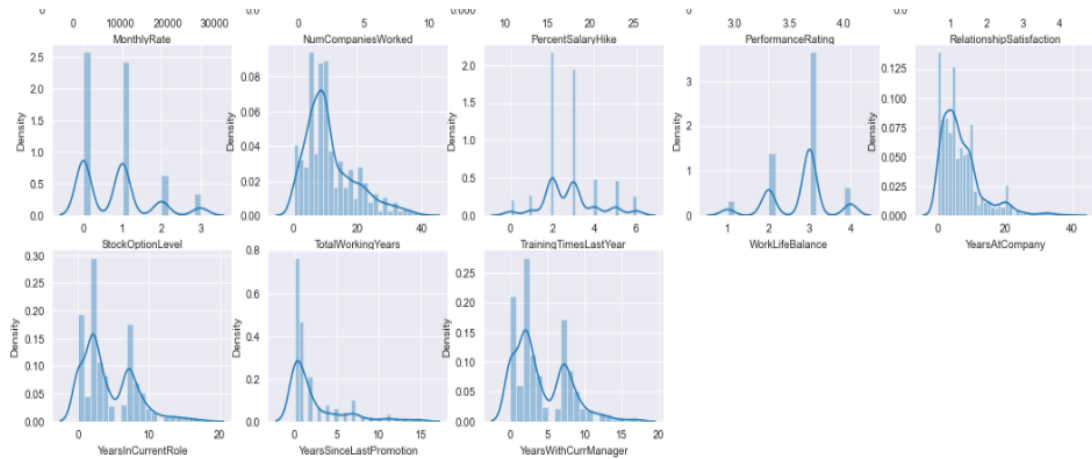


employee who has worked with 1 company previously are more looking for job change

## Data Distribution

```
In [45]: dist=df_1.select_dtypes(exclude='object')
plt.figure(figsize=(18,15))
plot=1
for col in dist:
    if plot<=25:
        plt.subplot(5,5,plot)
        sns.distplot(df_1[col])
        plt.xlabel(col)
        plot=plot+1
plt.show();
```



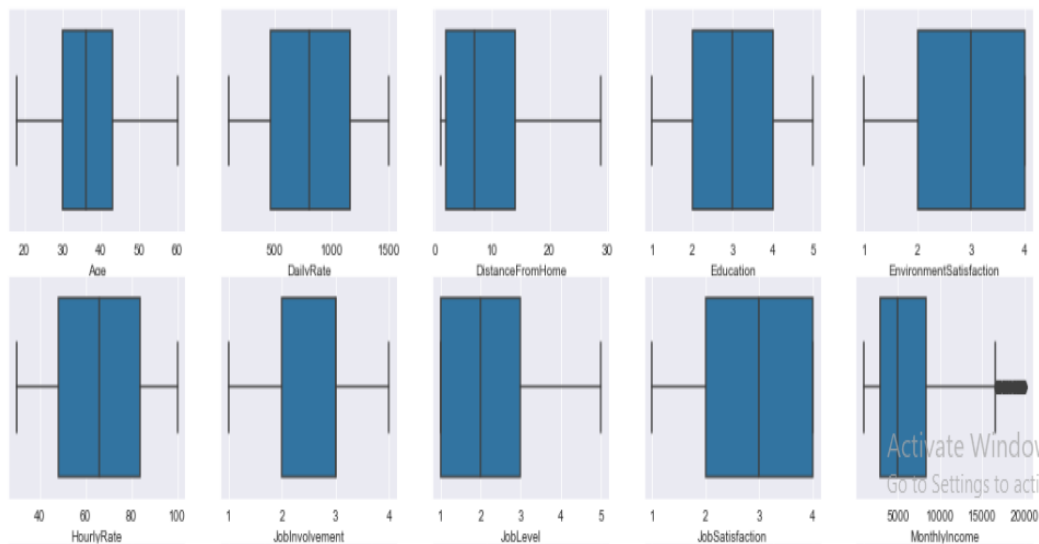


- 1- Most of our data is Normally distributed.
- 2- Attributes like Total working years, Years at company, Years since last promotion etc. are right skewed.
- 3- This shows some outliers must be present in our dataset.
- 4- we will remove skewness by some transformation methods.

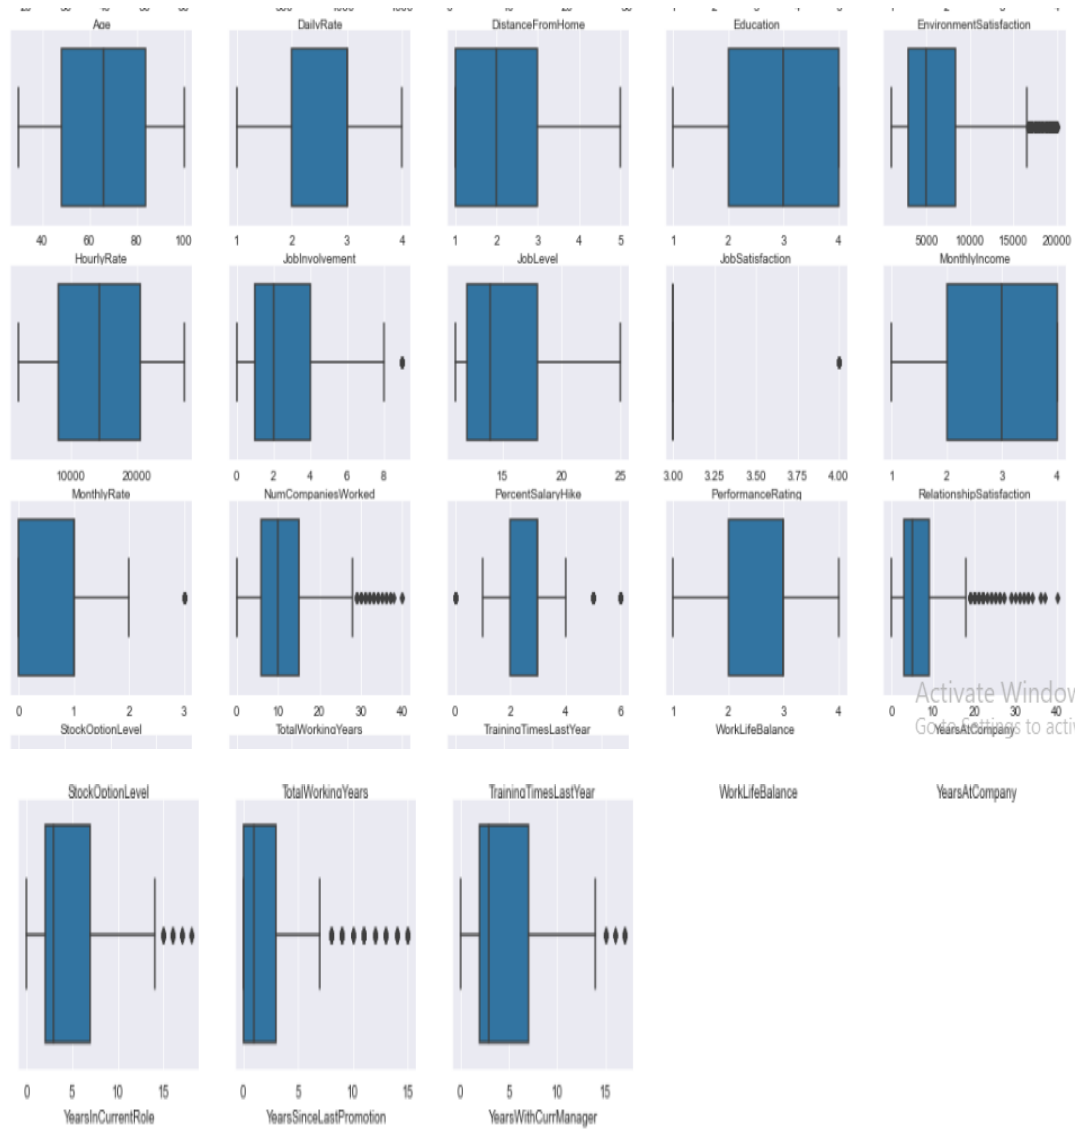
Activate Windows  
Go to Settings to activate Windows

## Finding outliers

```
In [46]: plt.figure(figsize=(18,15))
plot=1
for col in dist:
    if plot<=25:
        plt.subplot(5,5,plot)
        sns.boxplot(df_1[col])
        plt.xlabel(col)
        plot=plot+1
plt.show();
```



Activate Windows  
Go to Settings to activate Windows



```
In [47]: ## Removing Outliers
from scipy.stats import zscore
z = np.abs(zscore(dist))
print(z.shape)
df_1 = df_1.loc[(z<3).all(axis=1)]
print(df_1.shape)
```

(1470, 23)

(1387, 32)

## Skewness

```
In [48]: df_1.skew()
```

```
Out[48]: Age                0.472280  
DailyRate                -0.017078  
DistanceFromHome          0.954752  
Education                -0.289024  
EnvironmentSatisfaction  -0.325285  
HourlyRate               -0.030481  
JobInvolvement           -0.501401  
JobLevel                 1.126075  
JobSatisfaction          -0.345612  
MonthlyIncome            1.544770  
MonthlyRate              0.030596  
NumCompaniesWorked       1.037715  
PercentSalaryHike        0.800592  
PerformanceRating        1.931566  
RelationshipSatisfaction  -0.295686  
StockOptionLevel         0.962332  
TotalWorkingYears        1.034487  
TrainingTimesLastYear    0.577614  
WorkLifeBalance          -0.557100  
YearsAtCompany           1.248623  
YearsInCurrentRole       0.726675  
YearsSinceLastPromotion  1.756335  
YearsWithCurrManager     0.694506  
dtype: float64
```

Activate Windows  
Go to Settings to activate

```
In [49]: for i in dist:
         if df_1[i].skew() > .55:
             df_1[i] = np.log1p(df_1[i])
```

```
In [50]: df_1.skew()
```

```
Out[50]: Age                0.472280
         DailyRate          -0.017078
         DistanceFromHome   -0.031570
         Education          -0.289024
         EnvironmentSatisfaction -0.325285
         HourlyRate         -0.030481
         JobInvolvement     -0.501401
         JobLevel           0.497167
         JobSatisfaction    -0.345612
         MonthlyIncome       0.318873
         MonthlyRate        0.030596
         NumCompaniesWorked  0.101288
         PercentSalaryHike   0.496106
         PerformanceRating   1.931566
         RelationshipSatisfaction -0.295686
         StockOptionLevel    0.275912
         TotalWorkingYears   -0.728348
         TrainingTimesLastYear -1.044321
         WorkLifeBalance     -0.557100
         YearsAtCompany     -0.379527
         YearsInCurrentRole  -0.390406
         YearsSinceLastPromotion 0.695348
         YearsWithCurrManager -0.347018
         dtype: float64
```

Activate Windows  
Go to Settings to activate

## Encoding Categorical Value


```
In [51]: from sklearn.preprocessing import LabelEncoder
         le = LabelEncoder()
         for col in ob:
             df_1[col] = le.fit_transform(df_1[col])
```

```
In [52]: df_1.head()
```

```
Out[52]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	Gender	...	Performance
0	41	1	2	1102	2	0.693147	2	1	2	0	...	1.3
1	49	0	1	279	1	2.197225	1	1	3	1	...	1.6
2	37	1	2	1373	1	1.098612	2	4	4	1	...	1.3
3	33	0	1	1392	1	1.386294	4	1	4	0	...	1.3
4	27	0	2	591	1	1.098612	1	3	1	1	...	1.3

5 rows x 32 columns





## Splitting the data

```
In [53]: x=df_1.drop(['Attrition','Over18'],axis=1)
         y=df_1[['Attrition']]
```

```
In [54]: x
```

```
Out[54]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	Gender	HourlyRate	...	Perfor
0	41	2	1102	2	0.693147	2	1	2	0	94	...	
1	49	1	279	1	2.197225	1	1	3	1	61	...	
2	37	2	1373	1	1.098612	2	4	4	1	92	...	
3	33	1	1392	1	1.386294	4	1	4	0	56	...	
4	27	2	591	1	1.098612	1	3	1	1	40	...	
...	...	...	...	...	...	...	...	...	...	...	...	...
1465	36	1	884	1	3.178054	2	3	3	1	41	...	
1466	39	2	613	1	1.945910	1	3	4	1	42	...	
1467	27	2	155	1	1.609438	3	1	2	1	87	...	
1468	49	1	1023	2	1.098612	3	3	4	1	63	...	
1469	34	2	628	1	2.197225	3	3	2	1	82	...	

1387 rows x 30 columns

Activate Windows  
Go to Settings to activate

```
In [55]: y
```

```
Out[55]:
```

	Attrition
0	1
1	0
2	1
3	0
4	0
...	...
1465	0
1466	0
1467	0
1468	0
1469	0

1387 rows x 1 columns

Feature Scaling

```
In [56]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
x_scaled=ss.fit_transform(x)
x=pd.DataFrame(x_scaled,columns=x.columns)
x
```

```
Out[56]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	Gender	HourlyRate	...
0	0.536681	0.593126	0.734325	1.405373	-1.502086	-0.876177	-0.940815	-0.665328	-1.229911	1.388670	...
1	1.442111	-0.905354	-1.307769	-0.496337	0.253886	-1.853858	-0.940815	0.251978	0.813067	-0.239091	...
2	0.083966	0.593126	1.406752	-0.496337	-1.028716	-0.876177	1.305159	1.169285	0.813067	1.290017	...
3	-0.368749	-0.905354	1.453896	-0.496337	-0.692855	1.079185	-0.940815	1.169285	-1.229911	-0.485721	...
4	-1.047821	0.593126	-0.533609	-0.496337	-1.028716	-1.853858	0.556501	-1.582635	0.813067	-1.274939	...
...	...	...	...	...	...	...	...	...	...	...	...
1382	-0.029213	-0.905354	0.193406	-0.496337	1.398978	-0.876177	0.556501	0.251978	0.813067	-1.225613	...
1383	0.310324	0.593126	-0.479021	-0.496337	-0.039518	-1.853858	0.556501	1.169285	0.813067	-1.176286	...
1384	-1.047821	0.593126	-1.615447	-0.496337	-0.432340	0.101504	-0.940815	-0.665328	0.813067	1.043387	...
1385	1.442111	-0.905354	0.538304	1.405373	-1.028716	0.101504	0.556501	1.169285	0.813067	-0.140439	...
1386	-0.255570	0.593126	-0.441802	-0.496337	0.253886	0.101504	0.556501	-0.665328	0.813067	0.796757	...

1387 rows x 30 columns

Activate Windows  
Go to Settings to activate

## feature importance

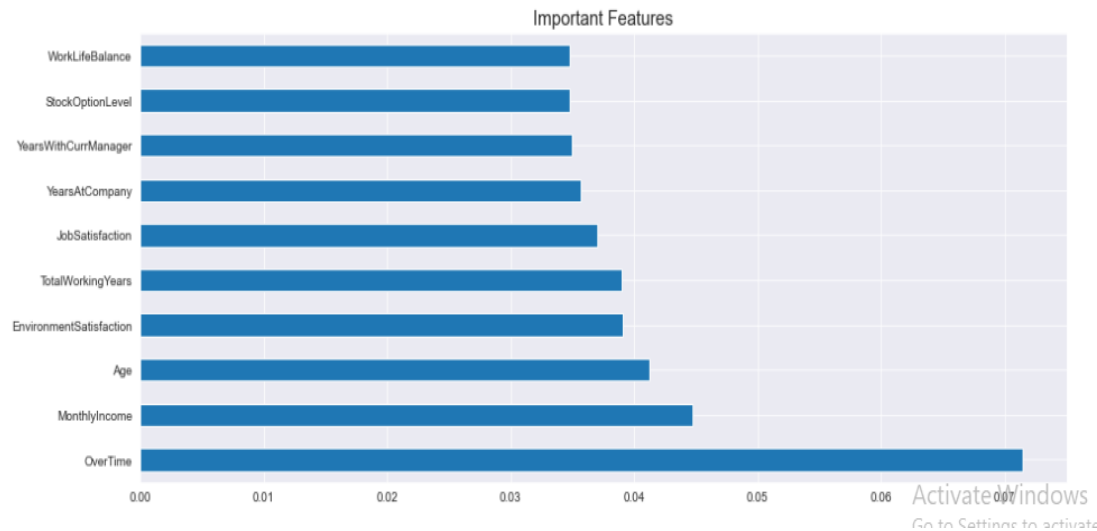
```
In [57]: from sklearn.ensemble import ExtraTreesClassifier
extra=ExtraTreesClassifier()
extra.fit(x,y)
```

```
Out[57]: ExtraTreesClassifier()
```

```
In [58]: print(extra.feature_importances_)
```

```
[0.04126763 0.02577878 0.03179439 0.02407701 0.03415031 0.02660997
 0.02816611 0.03906499 0.02080244 0.03137638 0.03272339 0.03367664
 0.03427738 0.03701433 0.03214459 0.0447477 0.032548 0.03174293
 0.07149409 0.0313795 0.01416903 0.03185999 0.03476039 0.03896788
 0.02932368 0.03475591 0.0357093 0.03246993 0.02813328 0.03501406]
```

```
In [59]: plt.figure(figsize=(15,6))
plt.title('Important Features', fontsize=15)
feat_importance=pd.Series(extra.feature_importances_, index=x.columns)
feat_importance.nlargest(10).plot(kind='barh')
plt.show();
```



## PCA

```
In [60]: from sklearn import decomposition
from sklearn.decomposition import PCA
covar_matrix=PCA(n_components=30)
```


```
In [61]: #Calculate Eigenvalues
covar_matrix.fit(x) ## x should be scaled
variance = covar_matrix.explained_variance_ratio_ #calculate variance ratios

var=np.cumsum(np.round(covar_matrix.explained_variance_ratio_, decimals=3)*100)
var #cumulative sum of variance explained with [n] features
## draw the graph
plt.ylabel('% Variance Explained')
plt.xlabel('# of Features')
plt.title('PCA Analysis')
plt.ylim(34,100.5)
plt.style.context('seaborn-whitegrid')

plt.plot(var)
```

Out[61]: [<matplotlib.lines.Line2D at 0x1ec02bc39d0>]

Out[61]: [



## Build the models

```
In [82]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_auc_score, f1_score, roc_curve, auc
```

```

In [83]: def max_accuracy_score(clf,x,y):
          max_accuracy=0
          for i in range(42,100):
              x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.20,random_state=i,stratify=y)
              clf.fit(x_train,y_train)
              pred=clf.predict(x_test)
              accuracy_check=accuracy_score(y_test,pred)
              if accuracy_check>max_accuracy:
                  max_accuracy=accuracy_check
                  final_r=i
          print('max accuracy score corresponding to',final_r,'is',max_accuracy*100)
          print('\n')
          print('cross validation score',cross_val_score(clf,x,y,scoring='accuracy').mean()*100)
          print('\n')
          print('Standard Deviation',cross_val_score(clf,x,y,scoring='accuracy').std()*100)
          print('\n')
          print('F1 score',f1_score(y_test,pred)*100)
          print('\n')
          print('Training accuracy',clf.score(x_train,y_train)*100)
          print('\n')
          print('Test Accuracy',clf.score(x_test,y_test)*100)
          print('\n')
          print('Confusion Matrix',confusion_matrix(y_test,pred))
          print('\n')
          print('Classification Report',classification_report(y_test,pred))
          print('\n')
          print('Roc_auc Score',roc_auc_score(y_test,pred)*100)
          return final_r

```

Activate Windows  
Go to Settings to activate

Classification Report		precision	recall	f1-score	support
0	0.90	0.97	0.93	232	
1	0.71	0.43	0.54	46	
accuracy			0.88	278	
macro avg		0.81	0.70	0.74	278
weighted avg		0.87	0.88	0.87	278

Roc\_auc Score 70.01499250374813

Out[84]: 44

## Decision Tree

```
In [85]: dt = DecisionTreeClassifier()  
max_accuracy_score(dt,x,y)
```

max accuracy score corresponding to 46 is 83.09352517985612

cross validation score 77.35942653819183

Standard Deviation 1.7987936464167595

F1 score 40.816326530612244

Training accuracy 100.0

Test Accuracy 79.13669064748201

Confusion Matrix [[200 32]  
[ 26 20]]

Test Accuracy 79.13669064748201

Confusion Matrix [[200 32]  
[ 26 20]]

Classification Report		precision	recall	f1-score	support
0	0.88	0.86	0.87	232	
1	0.38	0.43	0.41	46	
accuracy			0.79	278	
macro avg	0.63	0.65	0.64	278	
weighted avg	0.80	0.79	0.80	278	

Roc\_auc Score 64.84257871064467

Out[85]: 46

## KNN

```
In [86]: > knn = KNeighborsClassifier()  
max_accuracy_score(knn,x,y)
```

max accuracy score corresponding to 44 is 85.97122302158273

cross validation score 84.93208321429499

Standard Deviation 0.7225415674798252

F1 score 22.22222222222218

Training accuracy 87.2858431018936

Test Accuracy 84.89208633093526

Confusion Matrix [[230 2]  
[ 40 6]]

⏏ ⏪ ⏩ ⏴ ⏵ ⏶ ⏷

Confusion Matrix [[230 2]  
[ 40 6]]

Classification Report		precision	recall	f1-score	support
0	0.85	0.99	0.92	232	
1	0.75	0.13	0.22	46	
accuracy			0.85	278	
macro avg	0.80	0.56	0.57	278	
weighted avg	0.83	0.85	0.80	278	

Roc\_auc Score 56.09070464767616

Out[86]: 44

## Random Forest

```
In [87]: rf =RandomForestClassifier()  
max_accuracy_score(rf,x,y)
```

max accuracy score corresponding to 66 is 87.76978417266187

cross validation score 85.36399761057581

Standard Deviation 0.642139252890573

F1 score 28.07017543859649

Training accuracy 100.0

Test Accuracy 85.25179856115108

Confusion Matrix  $\begin{bmatrix} 229 & 3 \\ 38 & 8 \end{bmatrix}$

Confusion Matrix  $\begin{bmatrix} 229 & 3 \\ 38 & 8 \end{bmatrix}$

Classification Report			precision	recall	f1-score	support
0	0.86	0.99	0.92	232		
1	0.73	0.17	0.28	46		
accuracy			0.85	278		
macro avg	0.79	0.58	0.60	278		
weighted avg	0.84	0.85	0.81	278		

Roc\_auc Score 58.0491004497751

Out[87]: 66



## AdaBoost

```
In [88]: M adb= AdaBoostClassifier()  
max_accuracy_score(adb,x,y)
```

max accuracy score corresponding to 44 is 91.36690647482014

cross validation score 86.73402072565774

Standard Deviation 1.9947756669159884

F1 score 61.53846153846153

Training accuracy 89.54012623985572

Test Accuracy 89.20863309352518

Confusion Matrix [[224 8]  
[ 22 24]]

- -

Classification Report			precision	recall	f1-score	support
0	0.91	0.97	0.94	232		
1	0.75	0.52	0.62	46		
accuracy			0.89	278		
macro avg			0.83	0.74	0.78	278
weighted avg			0.88	0.89	0.88	278

Roc\_auc Score 74.36281859070466

Out[88]: 44

## Gradient Boost

```
In [89]: M gb = GradientBoostingClassifier()  
max_accuracy_score(gb,x,y)
```

max accuracy score corresponding to 60 is 89.568345323741

cross validation score 86.08524011116016

Standard Deviation 1.395928216614923

F1 score 46.37681159420289

Training accuracy 96.12263300270514

Test Accuracy 86.6906474820144

Confusion Matrix [[225 7]  
[ 30 16]]

Classification Report			precision	recall	f1-score	support
0	0.88	0.97	0.92	232		
1	0.70	0.35	0.46	46		
accuracy			0.87	278		
macro avg	0.79	0.66	0.69	278		
weighted avg	0.85	0.87	0.85	278		

Roc\_auc Score 65.88268365817093

Out[89]: 60

As we can see that Logistic Regression gives the best accuracy as compared to others model and Roc\_auc Score is also good among all the models. so, Logistic Regression is the best model for this data set

## Hyperparameter Tuning

```
In [91]: M x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=60,test_size=.20,stratify=y)
        ## Upsampling
        lg=LogisticRegression()
        param={'penalty':['l2','l1'],'C': [.0001,.001,.01,1,10], 'solver':['liblinear','saga']}

        grid=GridSearchCV(estimator=lg,param_grid=param,scoring='accuracy',n_jobs=-1)

        grid.fit(x_train,y_train)

        grid.best_params_
```

```
Out[91]: {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}
```

```
In [93]: M lg_final=LogisticRegression(C=10,penalty='l2',solver='liblinear')
        lg_final.fit(x_train,y_train)
        pred=lg_final.predict(x_test)
        print('Final Accuracy_score :',accuracy_score(pred,y_test)*100)
        print('\n')
        print('Final f_1 score :',f1_score(pred,y_test)*100)
        print('\n')
        print('Final roc_auc score :',roc_auc_score(pred,y_test))
        print('\n')
        print('Final classification Report :',classification_report(pred,y_test))
        print('\n')
        print('Final confusion Matrix :',confusion_matrix(pred,y_test))
```

```
Final Accuracy_score : 89.92805755395683
```

```
Final f_1 score : 60.0
```

```
Final roc_auc score : 0.8882874015748031
```

```
Final classification Report :
```

			precision	recall	f1-score	support
	0	0.99	0.90	0.94		254
	1	0.46	0.88	0.60		24
	accuracy			0.90		278
	macro avg	0.72	0.89	0.77		278
	weighted avg	0.94	0.90	0.91		278

```
Final confusion Matrix : [[229 25]
 [ 3 21]]
```

Activate Windows  
Go to Settings to activate

## saving model

```
In [94]: import joblib
          joblib.dump(lg_final, 'Attrition_lg.pkl')
```

```
Out[94]: ['Attrition_lg.pkl']
```

### loading and testing

```
In [96]: loaded_model=joblib.load('Attrition_lg.pkl')
         prediction=loaded_model.predict(x_test)
```

```
In [97]: ▶ prediction
```

[illegible]